

Distributed Algorithms 2020

1a Introduction

Distributed problems

- **Initially:** each node only aware of *itself*

Distributed problems

- **Initially:** each node only aware of *itself*
- **Finally:** each node knows *its own part of the solution*

Distributed problems

- **Initially:** each node only aware of *itself*
- **Communication rounds:** each node exchanges messages with *its own neighbors*

Distributed problems

- **Initially:** each node only aware of *itself*
- **Communication rounds:** each node exchanges messages with *its own neighbors*
- **Finally:** each node knows *its own part of the solution*

Example: coloring

- Each node knows ***its own color***
- Why enough?
- Coloring used for scheduling:

"I know my color is 5"

=

"I can safely be active during time slot number 5"

Locality

Fast distributed algorithm

=

Small number of communication rounds

=

Nodes only use information
in their own **local neighborhood**

Tell everyone everything

- **Before round 1:** each node only knows about *itself*
- **During round 1:** everyone tells about themselves to their neighbors
- **After round 1:** everyone now knows about their *neighbors* — they see everything up to **distance 1**

Tell everyone everything

- **Before round 2:** everyone knows about their *neighbors* — i.e. up to **distance 1**
- **During round 2:** everyone can send this information to their neighbors
- **After round 2:** everyone also aware of their *neighbors' neighbors* — they see everything up to **distance 2**

Tell everyone everything

- **Before round 3:** everyone knows everything up to **distance 2**
- **During round 3:** everyone can send this information to their neighbors
- **After round 3:** everyone now knows everything up to **distance 3**

Tell everyone everything

- **Before round T :** everyone knows everything up to **distance $T - 1$**
- **During round T :** everyone can send this information to their neighbors
- **After round T :** everyone now knows everything up to **distance T**

Locality

- Which graph problems are **local**?
 - can be solved so that each node only looks at its own local neighborhood
- Which graph problems are **global**?
 - cannot be solved using only local information
 - we need to see the whole input graph

Two perspectives

- **Mathematics:**

- graph theory, locality, distances...

- **Communication networks:**

- computers, network connections, message-passing, algorithms...

Cost of communication

- **Communication:**

get *one* bit from another computer
in the same local network \approx *0.5 milliseconds*

- **Computation:**

one *billion* arithmetic operations \approx *0.5 milliseconds*

Understanding nature

- What are the **fundamental limitations** of all kinds of systems that consist of interacting entities?
 - computer networks
 - biological systems
 - social networks
 - job markets
 - animal populations ...

Negative results

- How to prove that some problems cannot be solved efficiently in distributed settings?
- How to prove that some problems are inherently global?